

CSC 665: Artificial Intelligence

Games: Adversarial Search

Instructor: Pooyan Fazli
San Francisco State University

How to Succeed in This Course

- Read all the slides
- Read all the required reading material
- Read the “Review Materials” on iLearn
- Concepts in this course take some time to sink in: be careful not to fall behind. Cramming will NOT do it
- Be active in lectures and on iLearn.
- Study in groups
- Ask us if you have the slightest doubt

Course Information: Feedback

- Please give feedback (positive or negative) as often as and as early as you can.

CSC 665: Anonymous Feedback

Name (Optional)

Email Address (Optional)

Any Feedback on CSC 665?

Submit

Never submit passwords through Google Forms.

<https://forms.gle/jeuK3BNmGW7vRaGW7>

Game Playing State-of-the-Art

■ Checkers:

- 1950: First computer player.
- 1959: Samuel's self-taught program.
- 1994: First computer world champion: Chinook ended 40-year-reign of human champion Marion Tinsley.
- 2007: Checkers solved! Endgame database of 39 trillion states


■ Chess:

- 1945-1960: Zuse, Wiener, Shannon, Turing, Newell & Simon, McCarthy.
- 1960s onward: gradual improvement
- 1997: special-purpose chess machine Deep Blue defeats human champion Gary Kasparov in a six-game match. Deep Blue examined 200M positions per second and extended some lines of search up to 40 ply. Current programs running on a PC rate > 3200 (vs 2870 for Magnus Carlsen).


■ Go:

- 1968: Zobrist's program plays legal Go, barely (b>200!)
- 2005-2014: Monte Carlo tree search enables rapid advances: current programs beat strong amateurs, and professionals
- 2016: Google AI Beat World Go Champion

■ Pacman



Behavior from Computation



Video of Demo Mystery Pacman



Adversarial Games


Types of Games

- Many different kinds of games!
- Axes:
 - Deterministic or stochastic?
 - Perfect information (fully observable)?
 - One, two, or more players?
 - Turn-taking or simultaneous?
 - Zero sum?
- Want algorithms for calculating a **strategy (policy)** which recommends a move from each state

Deterministic Games


- Many possible formalizations, one is:
 - States: S (start at s_0)
 - Players: $P=\{1\dots N\}$ (usually take turns)
 - Actions: A (may depend on player / state)
 - Transition Function: $S \times A \rightarrow S'$
 - Terminal Test: $S \rightarrow \{t, f\}$
 - Terminal Utilities: $S \times P \rightarrow R$
- Solution for a player is a **policy**: $S \rightarrow A$

Zero-Sum Games



- Zero-Sum Games

- Agents have *opposite* utilities
- Pure competition:
 - One *maximizes*, the other *minimizes*




- General Games

- Agents have *independent* utilities
- Cooperation, indifference, competition and more are all possible


Adversarial Search

Single-Agent Trees



Value of a State

Value of a state:
The best achievable
outcome (utility)
from that state



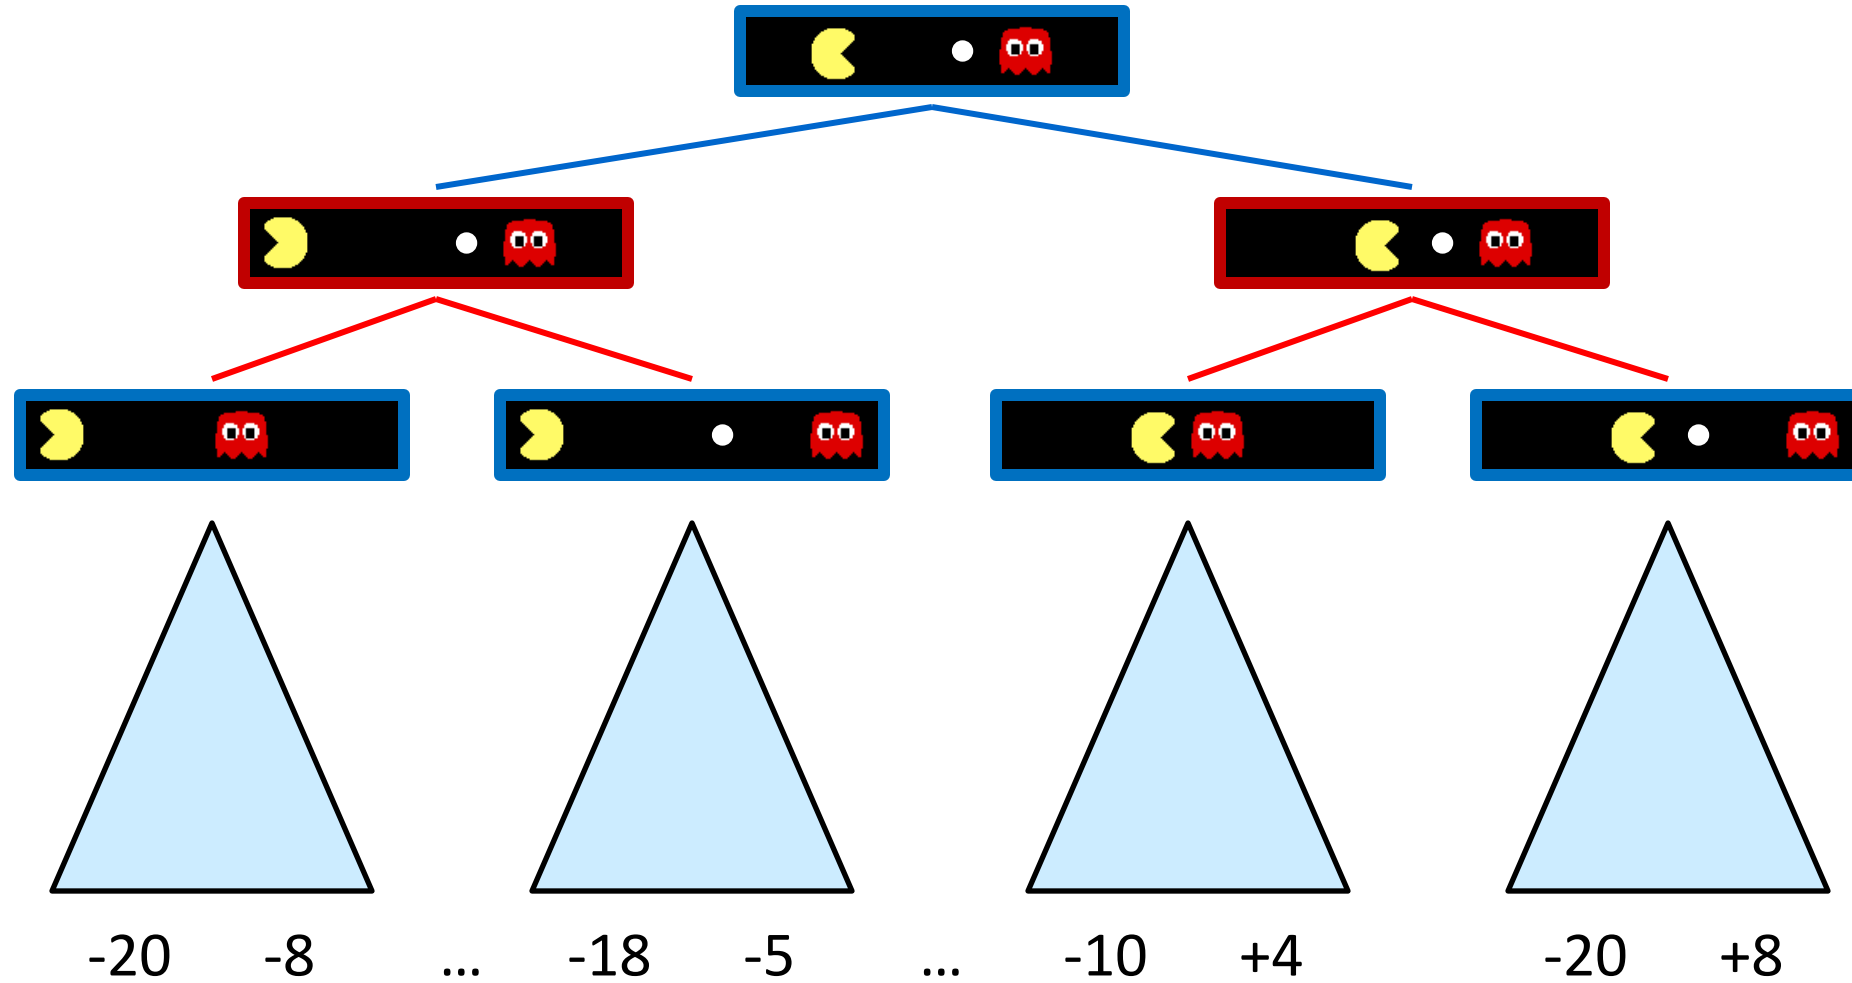
Non-Terminal States:

$$V(s) = \max_{s' \in \text{children}(s)} V(s')$$

Terminal States:

$$V(s) = \text{known}$$

Adversarial Game Trees




Minimax Values

States Under Agent's Control:

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$

States Under Opponent's Control:

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$



Terminal States:

$$V(s) = \text{known}$$

Tic-Tac-Toe Game Tree



MAX (X)




MIN (O)



MAX (X)



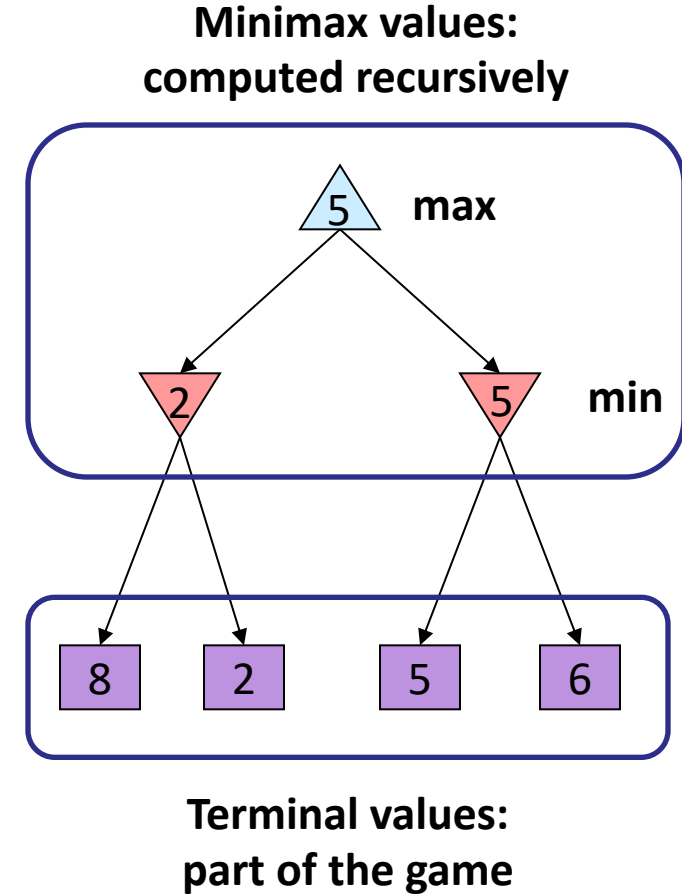
MIN (O)



The game tree for tic-tac-toe has fewer than 362,880 leaf nodes.

Adversarial Search (Minimax)

- **Deterministic, zero-sum games:**
 - Tic-tac-toe, chess, checkers
 - One player maximizes result
 - The other minimizes result
- **Minimax search:**
 - A state-space search tree
 - Players alternate turns
 - Compute each node's **minimax value**: the best achievable utility against a rational (optimal) adversary



Minimax Implementation

def max-value(state):

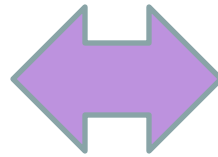
 initialize $v = -\infty$

 for each successor of state:

$v = \max(v, \text{min-value}(\text{successor}))$

 return v

$$V(s) = \max_{s' \in \text{successors}(s)} V(s')$$



def min-value(state):

 initialize $v = +\infty$

 for each successor of state:

$v = \min(v, \text{max-value}(\text{successor}))$

 return v

$$V(s') = \min_{s \in \text{successors}(s')} V(s)$$

Minimax Implementation (Dispatch)

```
def value(state):
```

if the state is a terminal state: return the state's utility

if the next agent is MAX: return max-value(state)

if the next agent is MIN: return min-value(state)

```
def max-value(state):
```

initialize $v = -\infty$

for each successor of state:

$v = \max(v, \text{min-value}(\text{successor}))$

return v


```
def min-value(state):
```

initialize $v = +\infty$

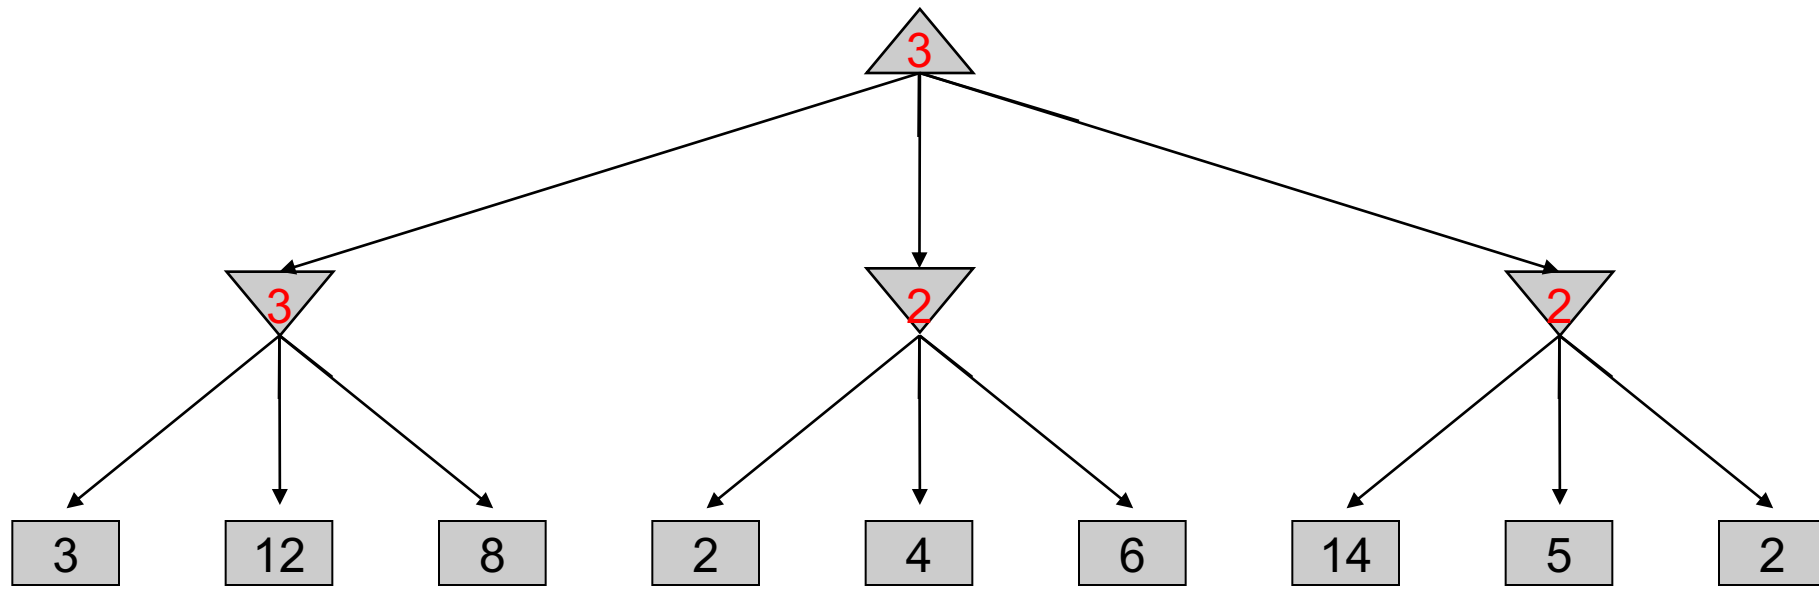
for each successor of state:

$v = \min(v, \text{max-value}(\text{successor}))$

return v



Minimax Example



Minimax Efficiency

- How efficient is minimax?
 - Just like (exhaustive) DFS
 - Time: $O(b^m)$
 - Space: $O(bm)$
- Example: For chess, $b \approx 35$, $m \approx 100$
 - Exact solution is completely infeasible
 - But, do we need to explore the whole tree?

Resource Limits

Resource Limits

- Problem: In realistic games, cannot search to leaves!

- Solution: Depth Limited Search


- Search only to a preset **depth limit** or **horizon**
- Use an **evaluation function** for non-terminal positions

- Guarantee of optimal play is gone


- Deeper search makes a BIG difference

- Example:


- Suppose we have 100 seconds, can explore 10K nodes / sec
- So can check 1M nodes per move
- For chess, $b \sim 35$ so reaches about depth 4 – not so good



Video of Demo Limited Depth (2)




Video of Demo Limited Depth (10)



Evaluation Functions

Evaluation Functions

- Evaluation functions score non-terminals in depth-limited search




- Ideal evaluation function: returns the actual minimax value of the position
- In practice: typically weighted linear sum of features:

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$



- e.g. $f_1(s) = (\text{num white queens} - \text{num black queens})$,
 $f_2(s) = (\text{num white pawns} - \text{num black pawns})$,

Evaluation for Pacman



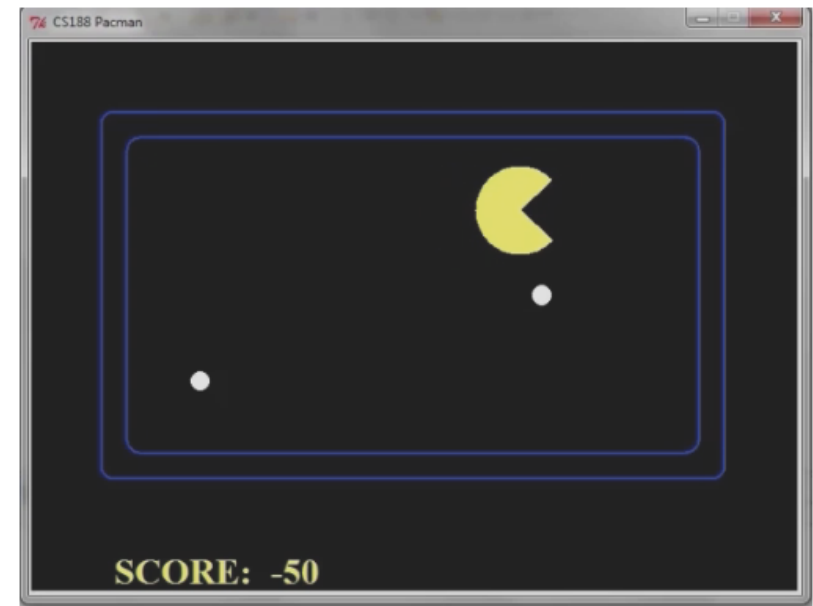


Evaluation Function Quiz

For the two situations shown below, which evaluation functions will give the situation on the left a higher score than the situation on the right?



1. $1 / (\text{Pac-Man's distance to the nearest food pellet})$
2. Pac-Man's distance to the nearest ghost ✓
3. Pac-Man's distance to the nearest ghost + $1 / (\text{Pac-Man's distance to the nearest food pellet})$ ✓
4. Pac-Man's distance to the nearest ghost + $1000 / (\text{Pac-Man's distance to the nearest food pellet})$


Video of Demo Thrashing (d=3)



Why Pacman Starves

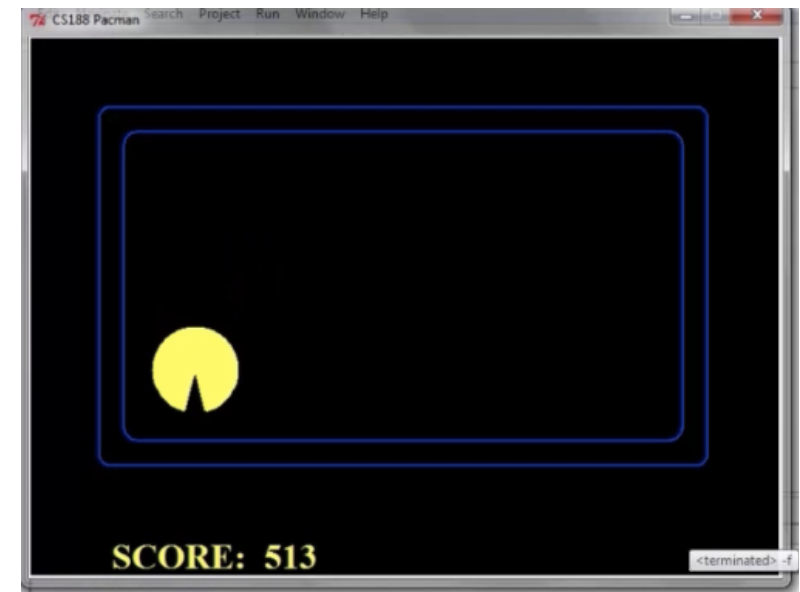


Eval Function:
Number of Dots Eaten

Better Eval Function:
Number of Dots Eaten +
 $1/\text{Closeness to the next dot}$







- A danger of replanning agents!
 - He knows his score will go up by eating the dot now
 - He knows his score will go up just as much by eating the dot later
 - Therefore, waiting seems just as good as eating.

Video of Demo Thrashing -- Fixed ($d=3$)




Video of Demo Smart Ghosts (Coordination)



Collaboration Quiz

Below is an example of a game tree with two minimizer players (min 1 and min 2), and one maximizer player.




What is the minimax value of this game tree? 6


Which action will the maximizer take when playing according to the minimax strategy? Left

Game Tree Pruning

Minimax Example




Minimax Pruning



Alpha-Beta Pruning

- General case (pruning children of MIN node)
 - We're computing the MIN-VALUE at some node n
 - We're looping over n 's children
 - n 's estimate of the children's min is dropping
 - Who cares about n 's value? MAX
 - Let α be the best value that MAX can get so far at any choice point along the current path from the root
 - If n becomes worse than α , MAX will avoid it, so we can prune n 's other children (it's already bad enough that it won't be played)
- Pruning children of MAX node is symmetric
 - Let β be the best value that MIN can get so far at any choice point along the current path from the root



Alpha-Beta Implementation

α : MAX's best option on path to root

β : MIN's best option on path to root

def max-value(state, α , β):

 initialize $v = -\infty$

 for each successor of state:

$v = \max(v, \text{value}(\text{successor}, \alpha, \beta))$

 if $v \geq \beta$

 return v

$\alpha = \max(\alpha, v)$

 return v

def min-value(state, α , β):

 initialize $v = +\infty$

 for each successor of state:

$v = \min(v, \text{value}(\text{successor}, \alpha, \beta))$

 if $v \leq \alpha$

 return v

$\beta = \min(\beta, v)$


 return v

Alpha-Beta Pruning Properties

- Theorem: This pruning has **no effect** on minimax value computed for the root!
- Good child ordering improves effectiveness of pruning
- With “perfect ordering”:
 - Time complexity drops to $O(b^{m/2})$
 - Doubles solvable depth!
 - 1M nodes/move => depth=8, respectable


Alpha-Beta Quiz

For the game tree shown below, which branches will be pruned by alpha-beta pruning? **f**



Alpha-Beta Quiz 2

For the game tree shown below, which branches will be pruned by alpha-beta pruning? **g and l**



Reading

- Chapter 5.1-5.3, 5.5 in the ALMA textbook