

CSC 665: Artificial Intelligence

Reinforcement Learning II

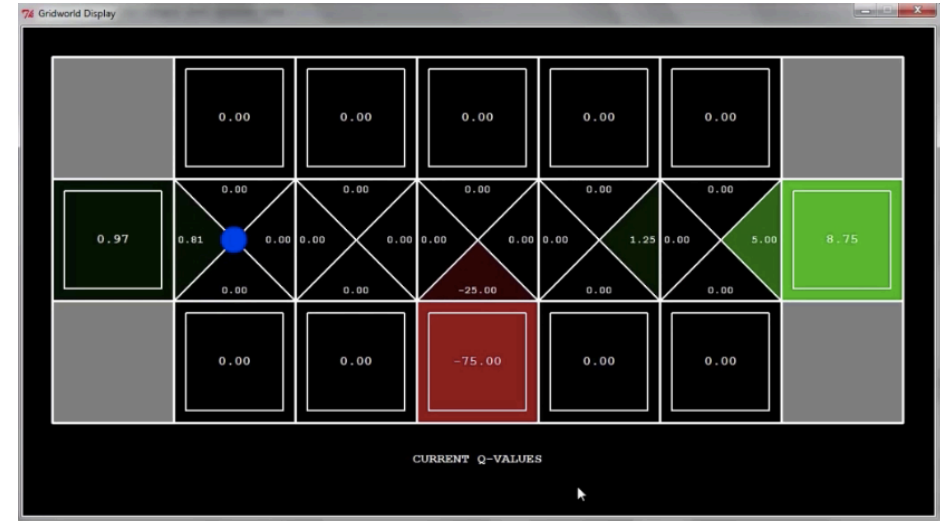
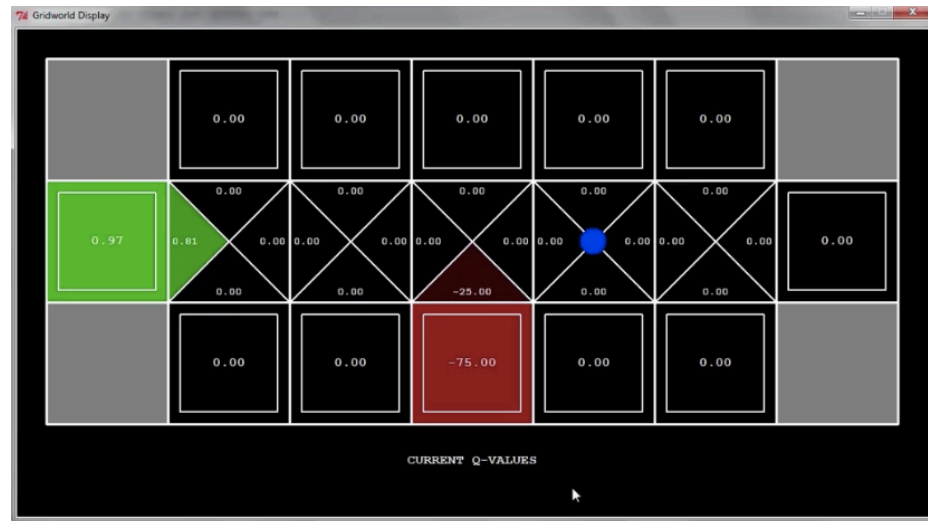
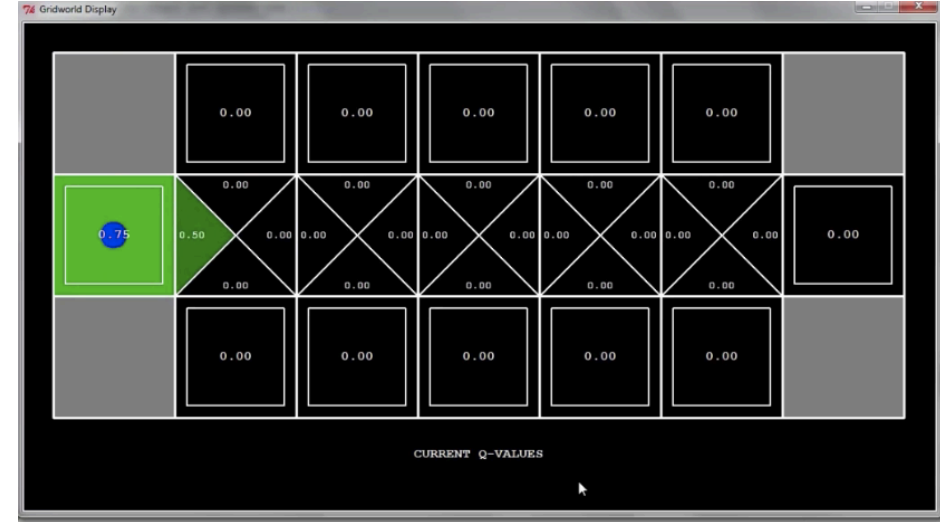
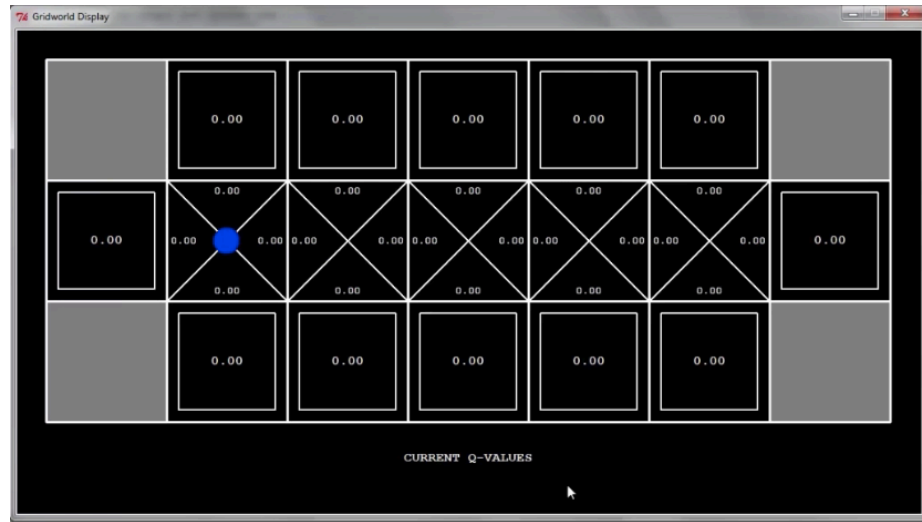
Instructor: Pooyan Fazli
San Francisco State University

Exploitation vs. Exploration

Exploitation vs. Exploration

- **Exploitation:** Execute the current best (maybe optimal?) policy to get high payoff
- **Exploration:** Try new sequences of (possibly random) actions to improve the agent's knowledge of the environment even though current model doesn't believe they have high payoff

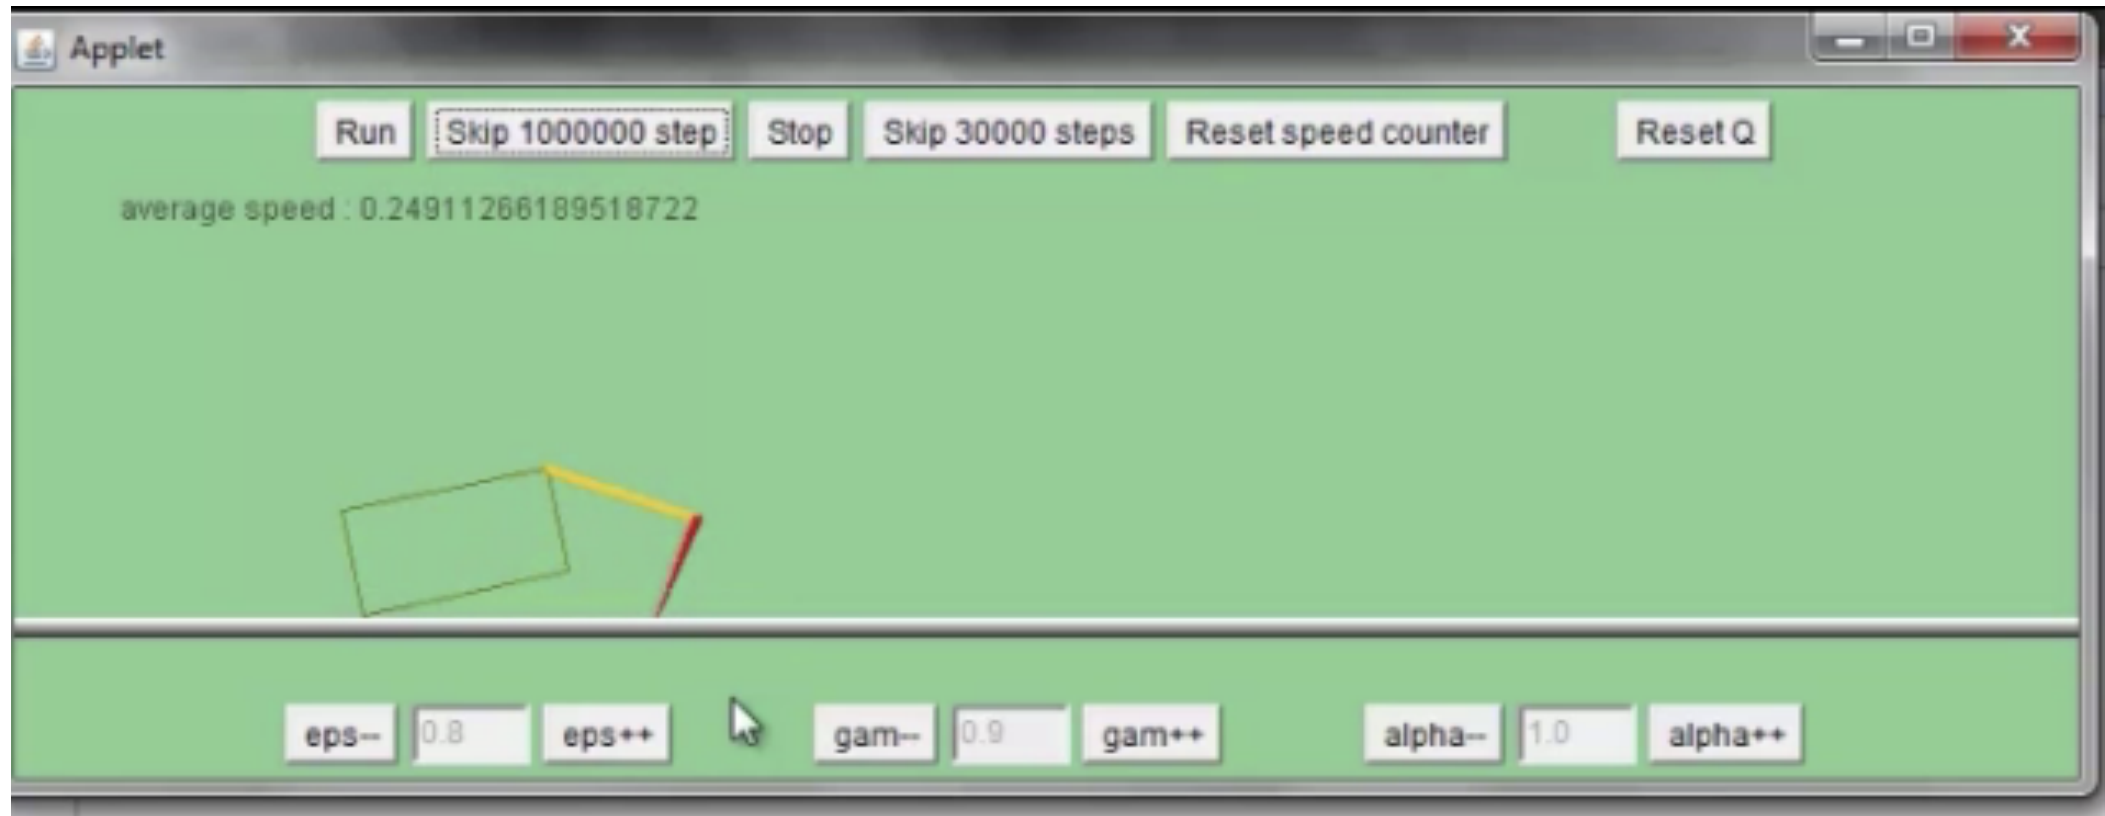
Q-learning – Manual Exploration – Bridge Grid



How to Explore?

- Simplest: random actions (ϵ -greedy)
 - Every time step, flip a coin
 - With (small) probability ϵ , act randomly
 - With (large) probability $1-\epsilon$, act on current best policy
- Properties of ϵ -greedy exploration
 - Every s,a pair is tried infinitely often
 - Does a lot of stupid things
 - Jumping off a cliff lots of times to make sure it hurts
 - Keeps doing stupid things for ever
 - Decay ϵ towards 0

Q-learning – Epsilon-Greedy – Crawler



Exploration Functions

- What to explore?

- Actions towards unexplored regions are encouraged (much faster than ϵ -greedy!)

- Exploration function

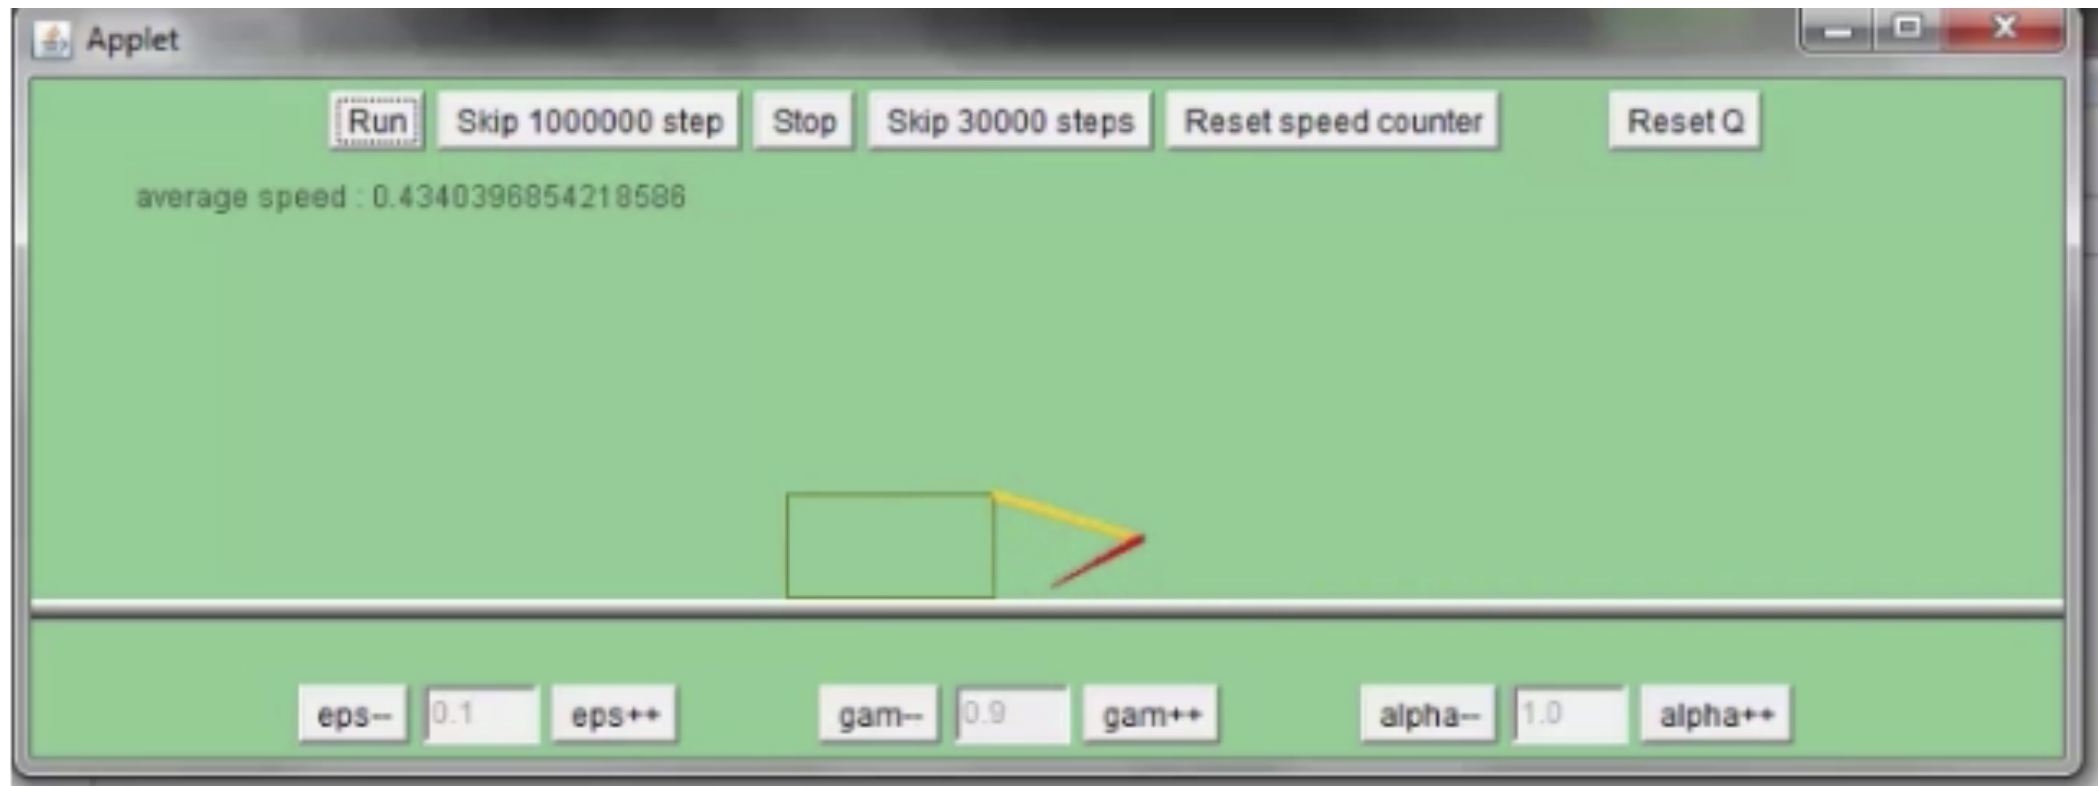
- Takes a value estimate u , a visit count n , and a constant k
returns:

$$f(u, n) = u + k/n$$

Regular Q-update: $Q(s,a) \leftarrow (1-\alpha) Q(s,a) + \alpha [R(s,a,s') + \gamma \max_{a'} Q(s',a')]]$

Modified Q-update: $Q(s,a) \leftarrow (1-\alpha) Q(s,a) + \alpha [R(s,a,s') + \gamma \max_{a'} f(Q(s',a'), n(s',a'))]$

Q-learning – Exploration Function – Crawler



Approximate Q-Learning

Generalizing Across States


- Basic Q-Learning keeps a table of all Q-values
- In realistic situations, we cannot possibly learn about every single state!
 - Too many states to visit them all in training
 - Too many states to hold the Q-tables in memory
- Instead, we want to generalize:
 - Learn about some small number of states from experience
 - Generalize that experience to new, similar situations

Example: Pacman


Let's say we discover through experience that this state is bad:



In naïve q-learning, we know nothing about this state:




Or even this one!



Feature-Based Representations

- **Solution:** describe a state using a vector of **features** (similar to the evaluation function in Assignment 2)
- **Features** are functions, $f(s)$, from states to real numbers that capture important properties of the state
 - Example features:
 - Distance to closest ghost
 - $1 / (\text{distance to closest ghost})$ f_{GST}
 - Number of ghosts
 - Distance to closest dot
 - $1 / (\text{dist to closest dot})$ f_{DOT}
 - Is Pacman in a tunnel? (0/1)
 - etc.
- Can also describe a q-state (s, a) with features (e.g. action a in state s moves closer to food)



Linear Value Functions

- Using a feature representation, we can write the Q-value (or the V-value) for any state using a few weights:

$$V(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

What you
used in
Assignment 2

$$Q(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \dots + w_n f_n(s, a)$$

What you use
in
Approximate
Q-learning

Approximate Q-Learning

$$Q(s, a) = w_1 f_1(s, a) + w_2 f_2(s, a) + \dots + w_n f_n(s, a)$$

- Q-learning with linear functions:


$$\text{transition} = (s, a, r, s')$$

$$\text{difference} = \left[r + \gamma \max_{a'} Q(s', a') \right] - Q(s, a)$$

$$Q(s, a) \leftarrow (1 - \alpha) Q(s, a) + (\alpha) \left[r + \gamma \max_{a'} Q(s', a') \right]$$

$$Q(s, a) \leftarrow Q(s, a) + \alpha [\text{difference}]$$

Q-Learning



What
you saw
before

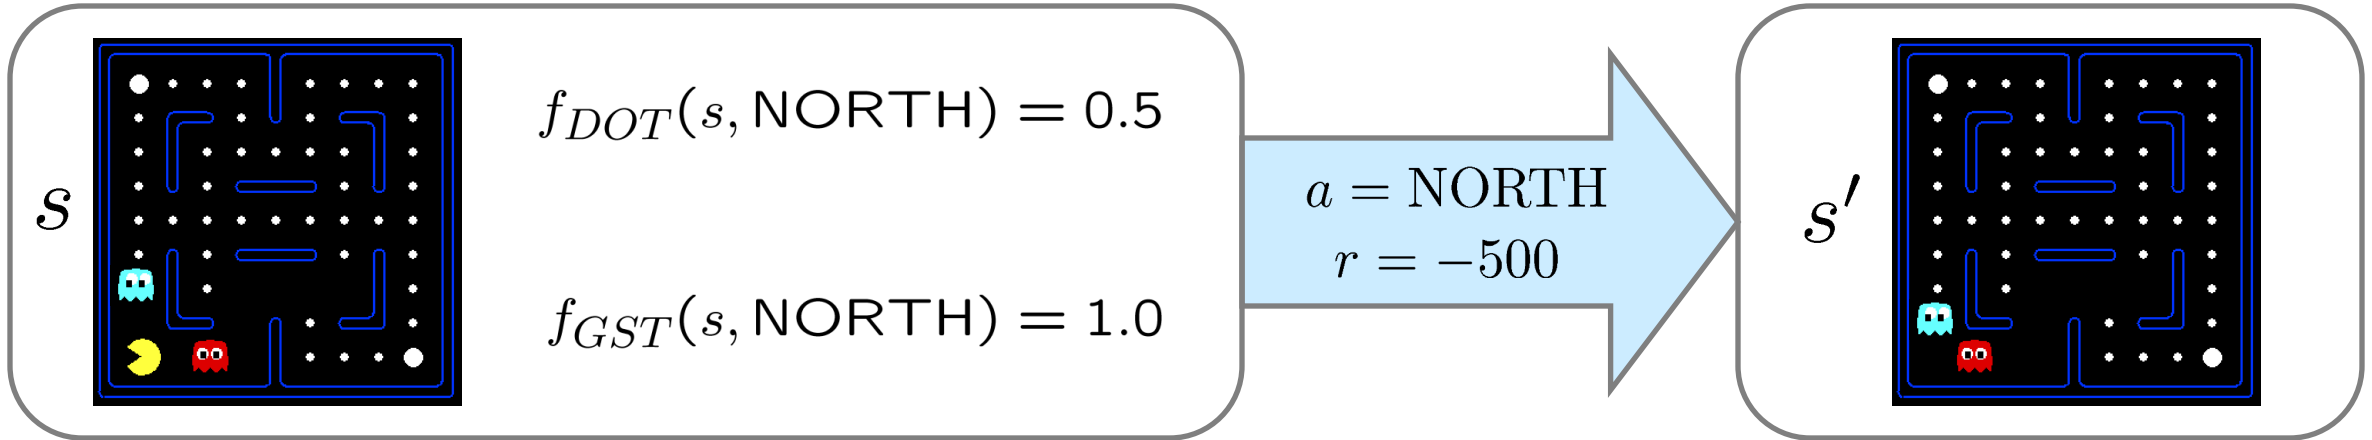
$$w_i \leftarrow w_i + \alpha [\text{difference}] f_i(s, a)$$

Approximate Q-Learning

- Intuitive interpretation:
 - Adjust weights of features

Example: Q-Pacman

$$Q(s, a) = 4.0 f_{DOT}(s, a) - 1.0 f_{GST}(s, a)$$



$$Q(s, \text{NORTH}) = +1$$

$$r + \gamma \max_{a'} Q(s', a') = -500 + 0$$

difference = -501 \longrightarrow




$$w_{DOT} \leftarrow 4.0 + \alpha [-501] 0.5$$

$$w_{GST} \leftarrow -1.0 + \alpha [-501] 1.0$$

$$Q(s, a) = 3.0 f_{DOT}(s, a) - 3.0 f_{GST}(s, a)$$

alpha = 0.004

Approximate Q-Learning -- Pacman



Conclusion

- We've seen how AI methods can solve problems in:
 - Search
 - Constraint Satisfaction Problems
 - Games
 - Markov Decision Problems
 - Reinforcement Learning
- Next up: Uncertainty and Learning!

Reading

- Read Section 22.4 in the ALMA textbook